

Chung Peng (Anderson) Lee

Email: cl6486@princeton.edu

URL: <https://anderson-lee-git.github.io/>

RESEARCH INTEREST

*I make dataset to assess properties of interest and make those in-distribution to improve models.
Currently thinking about physical safety / risk of multimodal LLMs, agents, and robotics.*

EDUCATION

- 2025 - now **Princeton University**
Ph.D. in Computer Science *Adviser: Aleksandra Korolova*
• **Relevant Courses** – Efficient Systems for Foundation Models, Deep Learning Theory, Reinforcement Learning, Information Theory, Safety-Critical Robotics
- 2021 - 2024 **University of Washington** GPA: 3.98 / 4.00
B.S. in Computer Science *Adviser: Jamie Morgenstern*
• Magna Cum Laude, Departmental Honor
• **Relevant Courses** – Machine Learning, Advanced Machine Learning, Deep Learning, Data Structures & Parallelism, Algorithms, Data Management (SQL), Explainable AI, Probabilistic Graphical Models, Natural Language Processing, Operating System

PUBLICATIONS

PREPRINTS

- 2026 **Chung Peng Lee**, Ryo Hachiuma, Hayoung Jung, Yusuke Hirota[†], and Aleksandra Korolova[†] (2026), “M3Risk: Assessing the Helpfulness of MLLMs for Malicious Intent-Driven Misuse,” In Submission
[†] denotes equal advising.
- Yida Yin*, Harish Krishnakumar*, **Chung Peng Lee**, Boya Zeng, Wenhao Chai, Shengbang Tong, Wenhui Chen, Hu Xu, Xingyu Fu, Gabriel Herbert Sarch, Aleksandra Korolova, and Zhuang Liu (2026), “WorldBench: A Challenging and Visually Diverse Multimodal Benchmark,” arXiv Preprint [Paper]
* denotes equal contribution.

CONFERENCE PAPERS

- 2026 Max Springer, **Chung Peng Lee**, Blossom Metevier, Jane Castleman, Bohdan Turbal, Hayoung Jung, Zeyu Shen, and Aleksandra Korolova (2026), “The Geometry of Alignment Collapse: When Fine-Tuning Breaks Safety,” The Fortieth Annual Conference on Neural Information Processing Systems (*NeurIPS 2026*). [Paper]
- Yung-Chin Chen, **Chung Peng Lee**, Ze-Wei Liou, and Naveen Verma (2026), “Massive Spikes in LLMs are Bias Vectors: Mechanistic Uncovering and Spike-Free Quantization,” The 2026 Conference on Empirical Methods in Natural Language Processing (*EMNLP 2026 Findings*). [Paper]
- Chung Peng Lee**, Rachel Hong, Harry H. Jiang, Aster Plotnik, William Agnew, and Jamie Morgenstern (2026), “How Do Data Owners Say No? A Case Study of Data Consent Mechanisms in Web-Scraped Vision-Language AI Training Datasets,” The 40th Annual AAAI Conference on Artificial Intelligence – AI for Societal Impact Track (*AISI 2026*). (Oral) [Paper]

- 2025 Yiwei Yang*, **Chung Peng Lee***, Shangbin Feng, Dora Zhao, Bingbing Wen, Anthony Z. Liu, Yulia Tsvetkov, and Bill Howe (2025), “Escaping the SpuriVerse: Can Large Vision-Language Models Generalize Beyond Seen Spurious Correlations?” The Thirty-Ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track (*NeurIPS D&B 2025*). [\[Paper\]](#)
 * denotes equal contribution.

WORKSHOP PAPERS

- 2024 **Chung Peng Lee**, Rachel Hong, and Jamie Morgenstern (2024), “Fairness Through Partial Awareness: Evaluation of the Addition of Demographic Information for Bias Mitigation Methods,” The Forty-First International Conference on Machine Learning – Next Generation of AI Safety Workshop (*ICML NGAIS 2024*). [\[Paper\]](#)

WORK EXPERIENCE

- 06/2026 **JPMorganChase** | *AI Research Intern* New York City, NY
 - 08/2026 *LLM Agent, Evaluation*
 • Black-Box Agent Adaptation
- 06/2024 **Stripe** | *Software Engineer (Machine Learning) Intern* Seattle, WA
 - 09/2024 *Airflow, Flyte, Chronon, Hadoop, (Py)Spark, Databricks, Bazel, AWS S3*
 • Migrated ML features to Chronon feature store: 2-day to <1-day latency, +\$0.6M/yr projected loss savings
- 06/2023 **SeekOut** | *Software Engineer Intern* Bellevue, WA
 - 08/2023 *Azure, Azure OpenAI, CosmosDB, NoSQL, Rest API, Kusto, C#, TypeScript, React*
 • Shipped an admin tool for client hiring managers to curate skills on their internal job boards, with skill recommendations

TEACHING

- 2026 COS 324: Introduction to Machine Learning (Teaching Assistant, Fall 2026)
 2024 CSE 446/546: Machine Learning (Teaching Assistant, Winter 2024)
 2023 CSE 446/546: Machine Learning (Teaching Assistant, Fall 2023)
 CSE 446/546: Machine Learning (Teaching Assistant, Spring 2023)
 CSE 123: Introductory Programming III (Teaching Assistant, Winter 2023)
 2022 CSE 121: Introductory Programming I (Teaching Assistant, Fall 2022)

SERVICE

- 2026 *Reviewer*, ICML 2026
 2025 *Program Committee*, AAAI – AI for Societal Impact Track (AISIT 2026)
Reviewer, NeurIPS Regulatable ML Workshop (NeurIPS RegML 2025)
 2024 *Reviewer*, ICML Next Generation of AI Safety Workshop (ICML NGAIS 2024)

Last updated: October 8, 2026